

```
# Load Pyspark Pkgs
import pyspark
```

```
# Methods/Attrib
dir(pyspark)
```

```
[ 'Accumulator',
  'AccumulatorParam',
  'Any',
  'BarrierTaskContext',
  'BarrierTaskInfo',
  'BasicProfiler',
  'Broadcast',
  'CPickleSerializer',
  'Callable',
  'HiveContext',
  'InheritableThread',
  'MarshalSerializer',
  'Optional',
  'Profiler',
  'RDD',
  'RDDBarrier',
  'Row',
  'SQLContext',
  'SparkConf',
  'SparkContext',
  'SparkFiles',
  'SparkJobInfo',
  'SparkStageInfo',
  'StatusTracker',
  'StorageLevel',
  'TaskContext',
  'TypeVar',
  'Union',
  '_F',
  '_NoValue',
  '__all__',
  '__builtins__',
  '__cached__',
  '__doc__',
  '__file__',
  '__loader__',
  '__name__',
  '__package__',
  '__path__',
  '__spec__',
  '__version__',
  '__globals__',
  'accumulators',
  'broadcast',
  'cast',
  'cloudpickle',
  'conf',
  'context',
  'copy_func',
  'errors',
  'files',
  'filterwarnings',
  'find_spark_home',
  'inheritable_thread_target',
  'java_gateway',
  'join',
  'keyword_only',
  'profiler',
```

✓ Working with DataFrames in PySpark

- Read DataSet(CSV)
- Create DataFrame

Tips

- SparkSession
- SparkContext :sc
- SqlContext

```
# Create A SparkSession
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.appName("PySparkTut").getOrCreate()
```

```
!ls
```

```
sample_data
```

```
!wget https://raw.githubusercontent.com/Jcharis/common_ml_datasets_explorer_app/master/datasets/diamonds.csv
```

```
--2024-11-12 07:00:33-- https://raw.githubusercontent.com/Jcharis/common_ml_datasets_explorer_app/master/datasets/diamonds.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 2772143 (2.6M) [text/plain]
Saving to: 'diamonds.csv'
```

```
diamonds.csv      100%[=====>]  2.64M  --.-KB/s   in 0.08s
```

```
2024-11-12 07:00:34 (34.0 MB/s) - 'diamonds.csv' saved [2772143/2772143]
```

```
!ls
```

```
diamonds.csv sample_data
```

```
# Read A DataSet without header
df = spark.read.csv('diamonds.csv')
```

```
# Preview dataset
df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|_c0|_c1|_c2|_c3|_c4|_c5|_c6|_c7|_c8|_c9|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|cut|color|clarity|depth|table|price|x|y|z|
|0.23|Ideal|E|SI2|61.5|55|326|3.95|3.98|2.43|
|0.21|Premium|E|SI1|59.8|61|326|3.89|3.84|2.31|
|0.23|Good|E|VS1|56.9|65|327|4.05|4.07|2.31|
|0.29|Premium|I|VS2|62.4|58|334|4.2|4.23|2.63|
|0.31|Good|J|SI2|63.3|58|335|4.34|4.35|2.75|
|0.24|Very Good|J|VVS2|62.8|57|336|3.94|3.96|2.48|
|0.24|Very Good|I|VVS1|62.3|57|336|3.95|3.98|2.47|
|0.26|Very Good|H|SI1|61.9|55|337|4.07|4.11|2.53|
|0.22|Fair|E|VS2|65.1|61|337|3.87|3.78|2.49|
|0.23|Very Good|H|VS1|59.4|61|338|4|4.05|2.39|
|0.3|Good|J|SI1|64|55|339|4.25|4.28|2.73|
|0.23|Ideal|J|VS1|62.8|56|340|3.93|3.9|2.46|
|0.22|Premium|F|SI1|60.4|61|342|3.88|3.84|2.33|
|0.31|Ideal|J|SI2|62.2|54|344|4.35|4.37|2.71|
|0.2|Premium|E|SI2|60.2|62|345|3.79|3.75|2.27|
|0.32|Premium|E|I1|60.9|58|345|4.38|4.42|2.68|
|0.3|Ideal|I|SI2|62|54|348|4.31|4.34|2.68|
|0.3|Good|J|SI1|63.4|54|351|4.23|4.29|2.7|
|0.3|Good|J|SI1|63.8|56|351|4.23|4.26|2.71|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
# Read A DataSet with header/column names
df = spark.read.csv('diamonds.csv',header=True)
```

```
df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|cut|color|clarity|depth|table|price|x|y|z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|0.23|Ideal|E|SI2|61.5|55|326|3.95|3.98|2.43|
|0.21|Premium|E|SI1|59.8|61|326|3.89|3.84|2.31|
|0.23|Good|E|VS1|56.9|65|327|4.05|4.07|2.31|
|0.29|Premium|I|VS2|62.4|58|334|4.2|4.23|2.63|
|0.31|Good|J|SI2|63.3|58|335|4.34|4.35|2.75|
|0.24|Very Good|J|VVS2|62.8|57|336|3.94|3.96|2.48|
|0.24|Very Good|I|VVS1|62.3|57|336|3.95|3.98|2.47|
|0.26|Very Good|H|SI1|61.9|55|337|4.07|4.11|2.53|
|0.22|Fair|E|VS2|65.1|61|337|3.87|3.78|2.49|
|0.23|Very Good|H|VS1|59.4|61|338|4|4.05|2.39|
|0.3|Good|J|SI1|64|55|339|4.25|4.28|2.73|
|0.23|Ideal|J|VS1|62.8|56|340|3.93|3.9|2.46|
```

```

| 0.22| Premium| F| SI1| 60.4| 61| 342|3.88|3.84|2.33|
| 0.31| Ideal| J| SI2| 62.2| 54| 344|4.35|4.37|2.71|
| 0.2| Premium| E| SI2| 60.2| 62| 345|3.79|3.75|2.27|
| 0.32| Premium| E| I1| 60.9| 58| 345|4.38|4.42|2.68|
| 0.3| Ideal| I| SI2| 62| 54| 348|4.31|4.34|2.68|
| 0.3| Good| J| SI1| 63.4| 54| 351|4.23|4.29| 2.7|
| 0.3| Good| J| SI1| 63.8| 56| 351|4.23|4.26|2.71|
| 0.3|Very Good| J| SI1| 62.7| 59| 351|4.21|4.27|2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

```
# Columns
```

```
df.columns
```

```
['carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z']
```

```
# Shape (rows + columns)
```

```
(df.count(), len(df.columns))
```

```
(53940, 10)
```

```
# Number of columns
```

```
len(df.columns)
```

```
10
```

```
# Number of rows
```

```
df.count()
```

```
53940
```

```
# Descriptive Analysis
```

```
df.describe().show()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|summary|      carat|      cut|color|clarity|      depth|      table|      price|      x|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| count|      53940|      53940|53940| 53940|      53940|      53940|      53940|      53940|
| mean|0.7979397478679852|      NULL| NULL| NULL| 61.74940489432624| 57.45718390804603|3932.799721913237| 5.731157211716609| 5.734525
| stddev|0.4740112444054196|      NULL| NULL| NULL| 1.4326213188336525| 2.2344905628213247|3989.439738146397| 1.1217607467924915| 1.1421346
| min|      0.2| Fair| D| I1|      43|      43|      1000|      0|
| max|      5.01|Very Good| J| VVS2|      79|      95|      9999|      9.86|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

```
# Pick a column & Get summary/describe a selected column
```

```
df.describe('carat').show()
```

```

+-----+-----+
|summary|      carat|
+-----+-----+
| count|      53940|
| mean|0.7979397478679852|
| stddev|0.4740112444054196|
| min|      0.2|
| max|      5.01|
+-----+-----+

```

```
# Preview the First Row
```

```
df.first()
```

```
Row(carat='0.23', cut='Ideal', color='E', clarity='SI2', depth='61.5', table='55', price='326', x='3.95', y='3.98', z='2.43')
```

```
# Preview the first 10 rows
```

```
# Like a list
```

```
df.head(10)
```

```

[Row(carat='0.23', cut='Ideal', color='E', clarity='SI2', depth='61.5', table='55', price='326', x='3.95', y='3.98', z='2.43'),
 Row(carat='0.21', cut='Premium', color='E', clarity='SI1', depth='59.8', table='61', price='326', x='3.89', y='3.84', z='2.31'),
 Row(carat='0.23', cut='Good', color='E', clarity='VS1', depth='56.9', table='65', price='327', x='4.05', y='4.07', z='2.31'),
 Row(carat='0.29', cut='Premium', color='I', clarity='VS2', depth='62.4', table='58', price='334', x='4.2', y='4.23', z='2.63'),
 Row(carat='0.31', cut='Good', color='J', clarity='SI2', depth='63.3', table='58', price='335', x='4.34', y='4.35', z='2.75'),

```

```
Row(carat='0.24', cut='Very Good', color='J', clarity='VVS2', depth='62.8', table='57', price='336', x='3.94', y='3.96', z='2.48'),
Row(carat='0.24', cut='Very Good', color='I', clarity='VVS1', depth='62.3', table='57', price='336', x='3.95', y='3.98', z='2.47'),
Row(carat='0.26', cut='Very Good', color='H', clarity='SI1', depth='61.9', table='55', price='337', x='4.07', y='4.11', z='2.53'),
Row(carat='0.22', cut='Fair', color='E', clarity='VS2', depth='65.1', table='61', price='337', x='3.87', y='3.78', z='2.49'),
Row(carat='0.23', cut='Very Good', color='H', clarity='VS1', depth='59.4', table='61', price='338', x='4', y='4.05', z='2.39')]
```

```
# Method 2: Useful Action with show()
# Show first 10 datapoints
df.show(10)
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|    cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|  Ideal|  E|   SI2| 61.5|  55| 326|3.95|3.98|2.43|
| 0.21| Premium|  E|   SI1| 59.8|  61| 326|3.89|3.84|2.31|
| 0.23|   Good|  E|   VS1| 56.9|  65| 327|4.05|4.07|2.31|
| 0.29| Premium|  I|   VS2| 62.4|  58| 334| 4.2|4.23|2.63|
| 0.31|   Good|  J|   SI2| 63.3|  58| 335|4.34|4.35|2.75|
| 0.24|Very Good|  J|  VVS2| 62.8|  57| 336|3.94|3.96|2.48|
| 0.24|Very Good|  I|  VVS1| 62.3|  57| 336|3.95|3.98|2.47|
| 0.26|Very Good|  H|   SI1| 61.9|  55| 337|4.07|4.11|2.53|
| 0.22|   Fair|  E|   VS2| 65.1|  61| 337|3.87|3.78|2.49|
| 0.23|Very Good|  H|   VS1| 59.4|  61| 338| 4|4.05|2.39|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 10 rows
```

```
# Get Last Rows
df.tail(5)
```

```

[Row(carat='0.72', cut='Ideal', color='D', clarity='SI1', depth='60.8', table='57', price='2757', x='5.75', y='5.76', z='3.5'),
Row(carat='0.72', cut='Good', color='D', clarity='SI1', depth='63.1', table='55', price='2757', x='5.69', y='5.75', z='3.61'),
Row(carat='0.7', cut='Very Good', color='D', clarity='SI1', depth='62.8', table='60', price='2757', x='5.66', y='5.68', z='3.56'),
Row(carat='0.86', cut='Premium', color='H', clarity='SI2', depth='61', table='58', price='2757', x='6.15', y='6.12', z='3.74'),
Row(carat='0.75', cut='Ideal', color='D', clarity='SI2', depth='62.2', table='55', price='2757', x='5.83', y='5.87', z='3.64')]
```

✓ Selection of columns

- .select

Note

- Dot & Bracket Notation only gives the column name not the entire column
 - [colA]*
 - .colA*

```
# List all Columns
df.columns
```

```
['carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z']
```

```
# Select A Column
df.select('carat').show()
```

```

+-----+
|carat|
+-----+
| 0.23|
| 0.21|
| 0.23|
| 0.29|
| 0.31|
| 0.24|
| 0.24|
| 0.26|
| 0.22|
| 0.23|
| 0.3|
| 0.23|
| 0.22|
| 0.31|
| 0.2|
| 0.32|
| 0.3|
| 0.3|
| 0.3|
```

```
| 0.3|
+-----+
only showing top 20 rows
```

```
# Select A Column irrespective of column word case
# will work irrespective of the case of the column once it is found within the dataset
df.select('CARAT').show()
```

```
↵ +-----+
  |CARAT|
  +-----+
  | 0.23|
  | 0.21|
  | 0.23|
  | 0.29|
  | 0.31|
  | 0.24|
  | 0.24|
  | 0.26|
  | 0.22|
  | 0.23|
  | 0.3 |
  | 0.23|
  | 0.22|
  | 0.31|
  | 0.2 |
  | 0.32|
  | 0.3 |
  | 0.3 |
  | 0.3 |
  | 0.3 |
  +-----+
only showing top 20 rows
```

```
# This is not as we would expect in pandas
# For Bracket Notation : pick column name not the entire column
df['carat']
```

```
↵ Column<'carat'>
```

```
# This is not as we would expect in pandas
# For Dot Notation : pick column name not the entire column
df.carat
```

```
↵ Column<'carat'>
```

```
# Select Multiple Columns
df.select('carat', 'cut').show(5)
```

```
↵ +-----+-----+
  |carat|    cut|
  +-----+-----+
  | 0.23| Ideal|
  | 0.21|Premium|
  | 0.23|  Good|
  | 0.29|Premium|
  | 0.31|  Good|
  +-----+-----+
only showing top 5 rows
```

▼ Column Filtering and Applying Conditions

- .filter
- .where

```
# Filter of Columns
# Apply A Condition
df.show(10)
```

```
↵ +-----+-----+-----+-----+-----+-----+-----+-----+
  |carat|    cut|color|clarity|depth|table|price|  x|  y|  z|
  +-----+-----+-----+-----+-----+-----+-----+-----+
  | 0.23|  Ideal|  E| SI2| 61.5|  55| 326|3.95|3.98|2.43|
```

```

| 0.21|Premium|E|SI1|59.8|61|326|3.89|3.84|2.31|
| 0.23|Good|E|VS1|56.9|65|327|4.05|4.07|2.31|
| 0.29|Premium|I|VS2|62.4|58|334|4.2|4.23|2.63|
| 0.31|Good|J|SI2|63.3|58|335|4.34|4.35|2.75|
| 0.24|Very Good|J|VVS2|62.8|57|336|3.94|3.96|2.48|
| 0.24|Very Good|I|VVS1|62.3|57|336|3.95|3.98|2.47|
| 0.26|Very Good|H|SI1|61.9|55|337|4.07|4.11|2.53|
| 0.22|Fair|E|VS2|65.1|61|337|3.87|3.78|2.49|
| 0.23|Very Good|H|VS1|59.4|61|338|4|4.05|2.39|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 10 rows

```

# Method 1:using filter
df.filter(df['cut'] == "Good").show()

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|cut|color|clarity|depth|table|price| x| y| z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|Good|E|VS1|56.9|65|327|4.05|4.07|2.31|
| 0.31|Good|J|SI2|63.3|58|335|4.34|4.35|2.75|
| 0.3|Good|J|SI1|64|55|339|4.25|4.28|2.73|
| 0.3|Good|J|SI1|63.4|54|351|4.23|4.29|2.7|
| 0.3|Good|J|SI1|63.8|56|351|4.23|4.26|2.71|
| 0.3|Good|I|SI2|63.3|56|351|4.26|4.3|2.71|
| 0.23|Good|F|VS1|58.2|59|402|4.06|4.08|2.37|
| 0.23|Good|E|VS1|64.1|59|402|3.83|3.85|2.46|
| 0.31|Good|H|SI1|64|54|402|4.29|4.31|2.75|
| 0.26|Good|D|VS2|65.2|56|403|3.99|4.02|2.61|
| 0.26|Good|D|VS1|58.4|63|403|4.19|4.24|2.46|
| 0.32|Good|H|SI2|63.1|56|403|4.34|4.37|2.75|
| 0.32|Good|H|SI2|63.8|56|403|4.36|4.38|2.79|
| 0.3|Good|I|SI1|63.2|55|405|4.25|4.29|2.7|
| 0.3|Good|H|SI1|63.7|57|554|4.28|4.26|2.72|
| 0.26|Good|E|VVS1|57.9|60|554|4.22|4.25|2.45|
| 0.7|Good|E|VS2|57.5|58|2759|5.85|5.9|3.38|
| 0.7|Good|F|VS1|59.4|62|2759|5.71|5.76|3.4|
| 0.7|Good|H|VVS2|62.1|64|2767|5.62|5.65|3.5|
| 0.71|Good|E|VS2|59.2|61|2772|5.8|5.88|3.46|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 20 rows

```

# Method 1:using filter
df.filter(df.carat >= 0.7).show()

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|cut|color|clarity|depth|table|price| x| y| z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.7|Ideal|E|SI1|62.5|57|2757|5.7|5.72|3.57|
| 0.86|Fair|E|SI2|55.1|69|2757|6.45|6.33|3.52|
| 0.7|Ideal|G|VS2|61.6|56|2757|5.7|5.67|3.5|
| 0.71|Very Good|E|VS2|62.4|57|2759|5.68|5.73|3.56|
| 0.78|Very Good|G|SI2|63.8|56|2759|5.81|5.85|3.72|
| 0.7|Good|E|VS2|57.5|58|2759|5.85|5.9|3.38|
| 0.7|Good|F|VS1|59.4|62|2759|5.71|5.76|3.4|
| 0.96|Fair|F|SI2|66.3|62|2759|6.27|5.95|4.07|
| 0.73|Very Good|E|SI1|61.6|59|2760|5.77|5.78|3.56|
| 0.8|Premium|H|SI1|61.5|58|2760|5.97|5.93|3.66|
| 0.75|Very Good|D|SI1|63.2|56|2760|5.8|5.75|3.65|
| 0.75|Premium|E|SI1|59.9|54|2760|6|5.96|3.58|
| 0.74|Ideal|G|SI1|61.6|55|2760|5.8|5.85|3.59|
| 0.75|Premium|G|VS2|61.7|58|2760|5.85|5.79|3.59|
| 0.8|Ideal|I|VS1|62.9|56|2760|5.94|5.87|3.72|
| 0.75|Ideal|G|SI1|62.2|55|2760|5.87|5.8|3.63|
| 0.8|Premium|G|SI1|63|59|2760|5.9|5.81|3.69|
| 0.74|Ideal|I|VVS2|62.3|55|2761|5.77|5.81|3.61|
| 0.81|Ideal|F|SI2|58.8|57|2761|6.14|6.11|3.6|
| 0.8|Ideal|F|SI2|61.4|57|2761|5.96|6|3.67|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 20 rows

```

# Method 2: where
df.where(df['cut'] == 'Good').show()

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|cut|color|clarity|depth|table|price| x| y| z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|Good|E|VS1|56.9|65|327|4.05|4.07|2.31|
| 0.31|Good|J|SI2|63.3|58|335|4.34|4.35|2.75|

```

```

| 0.3|Good|J|SI1|64|55|339|4.25|4.28|2.73|
| 0.3|Good|J|SI1|63.4|54|351|4.23|4.29|2.7|
| 0.3|Good|J|SI1|63.8|56|351|4.23|4.26|2.71|
| 0.3|Good|I|SI2|63.3|56|351|4.26|4.3|2.71|
| 0.23|Good|F|VS1|58.2|59|402|4.06|4.08|2.37|
| 0.23|Good|E|VS1|64.1|59|402|3.83|3.85|2.46|
| 0.31|Good|H|SI1|64|54|402|4.29|4.31|2.75|
| 0.26|Good|D|VS2|65.2|56|403|3.99|4.02|2.61|
| 0.26|Good|D|VS1|58.4|63|403|4.19|4.24|2.46|
| 0.32|Good|H|SI2|63.1|56|403|4.34|4.37|2.75|
| 0.32|Good|H|SI2|63.8|56|403|4.36|4.38|2.79|
| 0.3|Good|I|SI1|63.2|55|405|4.25|4.29|2.7|
| 0.3|Good|H|SI1|63.7|57|554|4.28|4.26|2.72|
| 0.26|Good|E|VVS1|57.9|60|554|4.22|4.25|2.45|
| 0.7|Good|E|VS2|57.5|58|2759|5.85|5.9|3.38|
| 0.7|Good|F|VS1|59.4|62|2759|5.71|5.76|3.4|
| 0.7|Good|H|VVS2|62.1|64|2767|5.62|5.65|3.5|
| 0.71|Good|E|VS2|59.2|61|2772|5.8|5.88|3.46|
+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

```

# Method 2: where
# select certain columns
df.where(df['cut'] == 'Good').select('price','cut').show()

```

```

+-----+-----+
|price| cut|
+-----+-----+
| 327|Good|
| 335|Good|
| 339|Good|
| 351|Good|
| 351|Good|
| 351|Good|
| 402|Good|
| 402|Good|
| 402|Good|
| 403|Good|
| 403|Good|
| 403|Good|
| 403|Good|
| 405|Good|
| 554|Good|
| 554|Good|
| 2759|Good|
| 2759|Good|
| 2767|Good|
| 2772|Good|
+-----+-----+
only showing top 20 rows

```

```

# Unique Values
# df['cut'].unique()
df.select("cut").distinct().show()

```

```

+-----+
| cut|
+-----+
| Premium|
| Ideal|
| Good|
| Fair|
| Very Good|
+-----+

```

How to Add Columns & Delete/Drop Columns

- .withColumn()
- .drop()

```

# Add Columns
df.withColumn("carat10x",df['carat'] * 10).show()

```

```

+-----+-----+-----+-----+-----+-----+-----+
|carat| cut|color|clarity|depth|table|price| x| y| z| carat10x|
+-----+-----+-----+-----+-----+-----+-----+

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23| Ideal| E| SI2| 61.5| 55| 326| 3.95| 3.98| 2.43| 2.3000000000000003|
| 0.21| Premium| E| SI1| 59.8| 61| 326| 3.89| 3.84| 2.31| 2.1|
| 0.23| Good| E| VS1| 56.9| 65| 327| 4.05| 4.07| 2.31| 2.3000000000000003|
| 0.29| Premium| I| VS2| 62.4| 58| 334| 4.2| 4.23| 2.63| 2.9|
| 0.31| Good| J| SI2| 63.3| 58| 335| 4.34| 4.35| 2.75| 3.1|
| 0.24| Very Good| J| VVS2| 62.8| 57| 336| 3.94| 3.96| 2.48| 2.4|
| 0.24| Very Good| I| VVS1| 62.3| 57| 336| 3.95| 3.98| 2.47| 2.4|
| 0.26| Very Good| H| SI1| 61.9| 55| 337| 4.07| 4.11| 2.53| 2.6|
| 0.22| Fair| E| VS2| 65.1| 61| 337| 3.87| 3.78| 2.49| 2.2|
| 0.23| Very Good| H| VS1| 59.4| 61| 338| 4| 4.05| 2.39| 2.3000000000000003|
| 0.3| Good| J| SI1| 64| 55| 339| 4.25| 4.28| 2.73| 3.0|
| 0.23| Ideal| J| VS1| 62.8| 56| 340| 3.93| 3.9| 2.46| 2.3000000000000003|
| 0.22| Premium| F| SI1| 60.4| 61| 342| 3.88| 3.84| 2.33| 2.2|
| 0.31| Ideal| J| SI2| 62.2| 54| 344| 4.35| 4.37| 2.71| 3.1|
| 0.2| Premium| E| SI2| 60.2| 62| 345| 3.79| 3.75| 2.27| 2.0|
| 0.32| Premium| E| I1| 60.9| 58| 345| 4.38| 4.42| 2.68| 3.2|
| 0.3| Ideal| I| SI2| 62| 54| 348| 4.31| 4.34| 2.68| 3.0|
| 0.3| Good| J| SI1| 63.4| 54| 351| 4.23| 4.29| 2.7| 3.0|
| 0.3| Good| J| SI1| 63.8| 56| 351| 4.23| 4.26| 2.71| 3.0|
| 0.3| Very Good| J| SI1| 62.7| 59| 351| 4.21| 4.27| 2.66| 3.0|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 20 rows

df.show()

```

↗ +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat| cut|color|clarity|depth|table|price| x| y| z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23| Ideal| E| SI2| 61.5| 55| 326| 3.95| 3.98| 2.43|
| 0.21| Premium| E| SI1| 59.8| 61| 326| 3.89| 3.84| 2.31|
| 0.23| Good| E| VS1| 56.9| 65| 327| 4.05| 4.07| 2.31|
| 0.29| Premium| I| VS2| 62.4| 58| 334| 4.2| 4.23| 2.63|
| 0.31| Good| J| SI2| 63.3| 58| 335| 4.34| 4.35| 2.75|
| 0.24| Very Good| J| VVS2| 62.8| 57| 336| 3.94| 3.96| 2.48|
| 0.24| Very Good| I| VVS1| 62.3| 57| 336| 3.95| 3.98| 2.47|
| 0.26| Very Good| H| SI1| 61.9| 55| 337| 4.07| 4.11| 2.53|
| 0.22| Fair| E| VS2| 65.1| 61| 337| 3.87| 3.78| 2.49|
| 0.23| Very Good| H| VS1| 59.4| 61| 338| 4| 4.05| 2.39|
| 0.3| Good| J| SI1| 64| 55| 339| 4.25| 4.28| 2.73|
| 0.23| Ideal| J| VS1| 62.8| 56| 340| 3.93| 3.9| 2.46|
| 0.22| Premium| F| SI1| 60.4| 61| 342| 3.88| 3.84| 2.33|
| 0.31| Ideal| J| SI2| 62.2| 54| 344| 4.35| 4.37| 2.71|
| 0.2| Premium| E| SI2| 60.2| 62| 345| 3.79| 3.75| 2.27|
| 0.32| Premium| E| I1| 60.9| 58| 345| 4.38| 4.42| 2.68|
| 0.3| Ideal| I| SI2| 62| 54| 348| 4.31| 4.34| 2.68|
| 0.3| Good| J| SI1| 63.4| 54| 351| 4.23| 4.29| 2.7|
| 0.3| Good| J| SI1| 63.8| 56| 351| 4.23| 4.26| 2.71|
| 0.3| Very Good| J| SI1| 62.7| 59| 351| 4.21| 4.27| 2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 20 rows

df2 = df.withColumn("carat10x",df['carat'] * 10)

df2.show()

```

↗ +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat| cut|color|clarity|depth|table|price| x| y| z| carat10x|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23| Ideal| E| SI2| 61.5| 55| 326| 3.95| 3.98| 2.43| 2.3000000000000003|
| 0.21| Premium| E| SI1| 59.8| 61| 326| 3.89| 3.84| 2.31| 2.1|
| 0.23| Good| E| VS1| 56.9| 65| 327| 4.05| 4.07| 2.31| 2.3000000000000003|
| 0.29| Premium| I| VS2| 62.4| 58| 334| 4.2| 4.23| 2.63| 2.9|
| 0.31| Good| J| SI2| 63.3| 58| 335| 4.34| 4.35| 2.75| 3.1|
| 0.24| Very Good| J| VVS2| 62.8| 57| 336| 3.94| 3.96| 2.48| 2.4|
| 0.24| Very Good| I| VVS1| 62.3| 57| 336| 3.95| 3.98| 2.47| 2.4|
| 0.26| Very Good| H| SI1| 61.9| 55| 337| 4.07| 4.11| 2.53| 2.6|
| 0.22| Fair| E| VS2| 65.1| 61| 337| 3.87| 3.78| 2.49| 2.2|
| 0.23| Very Good| H| VS1| 59.4| 61| 338| 4| 4.05| 2.39| 2.3000000000000003|
| 0.3| Good| J| SI1| 64| 55| 339| 4.25| 4.28| 2.73| 3.0|
| 0.23| Ideal| J| VS1| 62.8| 56| 340| 3.93| 3.9| 2.46| 2.3000000000000003|
| 0.22| Premium| F| SI1| 60.4| 61| 342| 3.88| 3.84| 2.33| 2.2|
| 0.31| Ideal| J| SI2| 62.2| 54| 344| 4.35| 4.37| 2.71| 3.1|
| 0.2| Premium| E| SI2| 60.2| 62| 345| 3.79| 3.75| 2.27| 2.0|
| 0.32| Premium| E| I1| 60.9| 58| 345| 4.38| 4.42| 2.68| 3.2|
| 0.3| Ideal| I| SI2| 62| 54| 348| 4.31| 4.34| 2.68| 3.0|
| 0.3| Good| J| SI1| 63.4| 54| 351| 4.23| 4.29| 2.7| 3.0|
| 0.3| Good| J| SI1| 63.8| 56| 351| 4.23| 4.26| 2.71| 3.0|
| 0.3| Very Good| J| SI1| 62.7| 59| 351| 4.21| 4.27| 2.66| 3.0|
+-----+-----+-----+-----+-----+-----+-----+-----+

```



```
+-----+
only showing top 20 rows
```

```
# Delete/Drop A Column
df2.drop('carat10x').show()
```

```
↗ +-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+
| 0.23|   Ideal|   E|   SI2| 61.5|  55|  326|3.95|3.98|2.43|
| 0.21| Premium|   E|   SI1| 59.8|  61|  326|3.89|3.84|2.31|
| 0.23|    Good|   E|   VS1| 56.9|  65|  327|4.05|4.07|2.31|
| 0.29| Premium|   I|   VS2| 62.4|  58|  334| 4.2|4.23|2.63|
| 0.31|    Good|   J|   SI2| 63.3|  58|  335|4.34|4.35|2.75|
| 0.24|Very Good|   J|  VVS2| 62.8|  57|  336|3.94|3.96|2.48|
| 0.24|Very Good|   I|  VVS1| 62.3|  57|  336|3.95|3.98|2.47|
| 0.26|Very Good|   H|   SI1| 61.9|  55|  337|4.07|4.11|2.53|
| 0.22|    Fair|   E|   VS2| 65.1|  61|  337|3.87|3.78|2.49|
| 0.23|Very Good|   H|   VS1| 59.4|  61|  338| 4|4.05|2.39|
| 0.3|    Good|   J|   SI1| 64|  55|  339|4.25|4.28|2.73|
| 0.23|   Ideal|   J|   VS1| 62.8|  56|  340|3.93| 3.9|2.46|
| 0.22| Premium|   F|   SI1| 60.4|  61|  342|3.88|3.84|2.33|
| 0.31|   Ideal|   J|   SI2| 62.2|  54|  344|4.35|4.37|2.71|
| 0.2| Premium|   E|   SI2| 60.2|  62|  345|3.79|3.75|2.27|
| 0.32| Premium|   E|    I1| 60.9|  58|  345|4.38|4.42|2.68|
| 0.3|   Ideal|   I|   SI2| 62|  54|  348|4.31|4.34|2.68|
| 0.3|    Good|   J|   SI1| 63.4|  54|  351|4.23|4.29| 2.7|
| 0.3|    Good|   J|   SI1| 63.8|  56|  351|4.23|4.26|2.71|
| 0.3|Very Good|   J|   SI1| 62.7|  59|  351|4.21|4.27|2.66|
+-----+
only showing top 20 rows
```

▼ GroupBy

- value counts
- aggregate

```
# Value Counts
# df['cut'].value_counts()
# df.groupby('cut')[].size()
df.groupby('cut').count().show()
```

```
↗ +-----+
|      cut|count|
+-----+
|   Premium|13791|
|    Ideal|21551|
|    Good| 4906|
|    Fair| 1610|
|Very Good|12082|
+-----+
```

```
# More Groupby
df.groupby('price').mean().show()
```

```
↗ +-----+
|price|
+-----+
| 2904|
| 3210|
| 3414|
| 3606|
| 3959|
| 4032|
| 4821|
| 4937|
| 5325|
| 6194|
| 6240|
| 6613|
| 6731|
| 7273|
| 7711|
| 7762|
| 9009|
```

```
| 9030|
| 9586|
|10096|
+-----+
only showing top 20 rows
```

```
# Sum of A groupby
df.groupBy('price').sum().show()
```

```
↗ +-----+
|price|
+-----+
| 2904|
| 3210|
| 3414|
| 3606|
| 3959|
| 4032|
| 4821|
| 4937|
| 5325|
| 6194|
| 6240|
| 6613|
| 6731|
| 7273|
| 7711|
| 7762|
| 9009|
| 9030|
| 9586|
|10096|
+-----+
only showing top 20 rows
```

```
# # Aggregation
# df.groupBy('carat').agg('col':'sum')
```

```
df.show()
```

```
↗ +-----+-----+-----+-----+-----+-----+-----+-----+
|carat|      cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|   Ideal|   E|   SI2| 61.5|  55|  326|3.95|3.98|2.43|
| 0.21|  Premium|   E|   SI1| 59.8|  61|  326|3.89|3.84|2.31|
| 0.23|    Good|   E|   VS1| 56.9|  65|  327|4.05|4.07|2.31|
| 0.29|  Premium|   I|   VS2| 62.4|  58|  334| 4.2|4.23|2.63|
| 0.31|    Good|   J|   SI2| 63.3|  58|  335|4.34|4.35|2.75|
| 0.24| Very Good|   J|  VVS2| 62.8|  57|  336|3.94|3.96|2.48|
| 0.24| Very Good|   I|  VVS1| 62.3|  57|  336|3.95|3.98|2.47|
| 0.26| Very Good|   H|   SI1| 61.9|  55|  337|4.07|4.11|2.53|
| 0.22|    Fair|   E|   VS2| 65.1|  61|  337|3.87|3.78|2.49|
| 0.23| Very Good|   H|   VS1| 59.4|  61|  338| 4|4.05|2.39|
| 0.3|    Good|   J|   SI1|  64|  55|  339|4.25|4.28|2.73|
| 0.23|   Ideal|   J|   VS1| 62.8|  56|  340|3.93| 3.9|2.46|
| 0.22|  Premium|   F|   SI1| 60.4|  61|  342|3.88|3.84|2.33|
| 0.31|   Ideal|   J|   SI2| 62.2|  54|  344|4.35|4.37|2.71|
| 0.2|    Premium|   E|   SI2| 60.2|  62|  345|3.79|3.75|2.27|
| 0.32|  Premium|   E|    I1| 60.9|  58|  345|4.38|4.42|2.68|
| 0.3|   Ideal|   I|   SI2|  62|  54|  348|4.31|4.34|2.68|
| 0.3|    Good|   J|   SI1| 63.4|  54|  351|4.23|4.29| 2.7|
| 0.3|    Good|   J|   SI1| 63.8|  56|  351|4.23|4.26|2.71|
| 0.3| Very Good|   J|   SI1| 62.7|  59|  351|4.21|4.27|2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
df.columns
```

```
↗ ['carat', 'cut', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z']
```

```
# Rearrange Columns
df.select('carat', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z', 'cut').show()
```

```
↗ +-----+-----+-----+-----+-----+-----+-----+-----+
|carat|color|clarity|depth|table|price|  x|  y|  z|      cut|
+-----+-----+-----+-----+-----+-----+-----+-----+
```

0.23	E	SI2	61.5	55	326	3.95	3.98	2.43	Ideal
0.21	E	SI1	59.8	61	326	3.89	3.84	2.31	Premium
0.23	E	VS1	56.9	65	327	4.05	4.07	2.31	Good
0.29	I	VS2	62.4	58	334	4.2	4.23	2.63	Premium
0.31	J	SI2	63.3	58	335	4.34	4.35	2.75	Good
0.24	J	VVS2	62.8	57	336	3.94	3.96	2.48	Very Good
0.24	I	VVS1	62.3	57	336	3.95	3.98	2.47	Very Good
0.26	H	SI1	61.9	55	337	4.07	4.11	2.53	Very Good
0.22	E	VS2	65.1	61	337	3.87	3.78	2.49	Fair
0.23	H	VS1	59.4	61	338	4	4.05	2.39	Very Good
0.3	J	SI1	64	55	339	4.25	4.28	2.73	Good
0.23	J	VS1	62.8	56	340	3.93	3.9	2.46	Ideal
0.22	F	SI1	60.4	61	342	3.88	3.84	2.33	Premium
0.31	J	SI2	62.2	54	344	4.35	4.37	2.71	Ideal
0.2	E	SI2	60.2	62	345	3.79	3.75	2.27	Premium
0.32	E	I1	60.9	58	345	4.38	4.42	2.68	Premium
0.3	I	SI2	62	54	348	4.31	4.34	2.68	Ideal
0.3	J	SI1	63.4	54	351	4.23	4.29	2.7	Good
0.3	J	SI1	63.8	56	351	4.23	4.26	2.71	Good
0.3	J	SI1	62.7	59	351	4.21	4.27	2.66	Very Good

only showing top 20 rows

```
# Assign DF to a New DF
```

```
new_df = df.select('carat', 'color', 'clarity', 'depth', 'table', 'price', 'x', 'y', 'z', 'cut')
```

```
new_df
```

```
DataFrame[carat: string, color: string, clarity: string, depth: string, table: string, price: string, x: string, y: string, z: string, cut: string]
```

```
# Check Datatype
```

```
new_df.dtypes
```

```
[('carat', 'string'),
 ('color', 'string'),
 ('clarity', 'string'),
 ('depth', 'string'),
 ('table', 'string'),
 ('price', 'string'),
 ('x', 'string'),
 ('y', 'string'),
 ('z', 'string'),
 ('cut', 'string')]
```

```
# Check For the Schema
```

```
df.printSchema()
```

```
root
|-- carat: string (nullable = true)
|-- cut: string (nullable = true)
|-- color: string (nullable = true)
|-- clarity: string (nullable = true)
|-- depth: string (nullable = true)
|-- table: string (nullable = true)
|-- price: string (nullable = true)
|-- x: string (nullable = true)
|-- y: string (nullable = true)
|-- z: string (nullable = true)
```

```
# Check type of DF
```

```
type(df)
```

```
pyspark.sql.dataframe.DataFrame
def __init__(jdf: JavaObject, sql_ctx: Union['SQLContext', 'SparkSession'])
    Supports Spark Connect.

Examples
-----
A :class:`DataFrame` is equivalent to a relational table in Spark SQL,
and can be created using various functions in :class:`SparkSession`:
```

✓ Saving DataFrames as CSV,parquet etc

```
# Save
new_df.write.format('csv').option('header', 'true').save("diamond_clean.csv")
```

```
!!ls
```

```
↩ diamond_clean.csv  diamonds.csv  sample_data
```

```
# Save as parquet
new_df.write.format('parquet').save("diamond_clean.parquet")
```

✓ Making SQL Queries

- parse in the spark.SparkContext
- sqlContext

```
from pyspark.sql import SQLContext
```

```
dir(spark)
```

```
↩ ['Builder',
  '__annotations__',
  '__class__',
  '__delattr__',
  '__dict__',
  '__dir__',
  '__doc__',
  '__enter__',
  '__eq__',
  '__exit__',
  '__format__',
  '__ge__',
  '__getattr__',
  '__gt__',
  '__hash__',
  '__init__',
  '__init_subclass__',
  '__le__',
  '__lt__',
  '__module__',
  '__ne__',
  '__new__',
  '__reduce__',
  '__reduce_ex__',
  '__repr__',
  '__setattr__',
  '__sizeof__',
  '__str__',
  '__subclasshook__',
  '__weakref__',
  '_activeSession',
  '_convert_from_pandas',
  '_createFromLocal',
  '_createFromRDD',
  '_create_dataframe',
  '_create_from_pandas_with_arrow',
  '_create_shell_session',
  '_getActiveSessionOrCreate',
  '_get_numpy_record_dtype',
  '_inferSchema',
  '_inferSchemaFromList',
  '_instantiatedSession',
  '_jconf',
  '_jsc',
  '_jsparkSession',
  '_jvm',
  '_repr_html_',
  '_sc',
  'active',
  'addArtifact',
  'addArtifacts',
  'addTag',
  'builder',
  'catalog',
  'clearTags',
  'client',
  'conf',
```

```
# Create A Spark Context From the Spark Session
sc = spark.sparkContext
```

```
# Parse into the SQLContext
sqlContext = SQLContext(sc)
```

```
⚡ /usr/local/lib/python3.10/dist-packages/pyspark/sql/context.py:113: FutureWarning: Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate().warn()
```

```
# Register Current DataFrame As Temporal Table
df.registerTempTable("DiamondsTable")
```

```
⚡ /usr/local/lib/python3.10/dist-packages/pyspark/sql/dataframe.py:329: FutureWarning: Deprecated in 2.0, use createOrReplaceTempView instead. inst
warnings.warn("Deprecated in 2.0, use createOrReplaceTempView instead.", FutureWarning)
```

```
# Making Quries
sqlContext.sql('SELECT * FROM DiamondsTable').show()
```

```
⚡ +-----+-----+-----+-----+-----+-----+-----+-----+-----+
|carat|    cut|color|clarity|depth|table|price|  x|  y|  z|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 0.23|  Ideal|  E|  SI2| 61.5| 55| 326|3.95|3.98|2.43|
| 0.21| Premium|  E|  SI1| 59.8| 61| 326|3.89|3.84|2.31|
| 0.23|   Good|  E|  VS1| 56.9| 65| 327|4.05|4.07|2.31|
| 0.29| Premium|  I|  VS2| 62.4| 58| 334| 4.2|4.23|2.63|
| 0.31|   Good|  J|  SI2| 63.3| 58| 335|4.34|4.35|2.75|
| 0.24|Very Good|  J| VVS2| 62.8| 57| 336|3.94|3.96|2.48|
| 0.24|Very Good|  I| VVS1| 62.3| 57| 336|3.95|3.98|2.47|
| 0.26|Very Good|  H|  SI1| 61.9| 55| 337|4.07|4.11|2.53|
| 0.22|   Fair|  E|  VS2| 65.1| 61| 337|3.87|3.78|2.49|
| 0.23|Very Good|  H|  VS1| 59.4| 61| 338| 4|4.05|2.39|
| 0.3|   Good|  J|  SI1| 64| 55| 339|4.25|4.28|2.73|
| 0.23|  Ideal|  J|  VS1| 62.8| 56| 340|3.93| 3.9|2.46|
| 0.22| Premium|  F|  SI1| 60.4| 61| 342|3.88|3.84|2.33|
| 0.31|  Ideal|  J|  SI2| 62.2| 54| 344|4.35|4.37|2.71|
| 0.2| Premium|  E|  SI2| 60.2| 62| 345|3.79|3.75|2.27|
| 0.32| Premium|  E|  I1| 60.9| 58| 345|4.38|4.42|2.68|
| 0.3|  Ideal|  I|  SI2| 62| 54| 348|4.31|4.34|2.68|
| 0.3|   Good|  J|  SI1| 63.4| 54| 351|4.23|4.29| 2.7|
| 0.3|   Good|  J|  SI1| 63.8| 56| 351|4.23|4.26|2.71|
| 0.3|Very Good|  J|  SI1| 62.7| 59| 351|4.21|4.27|2.66|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
# Can also use it to work with DataFrames
dir(sqlContext)
```

```
⚡ ['__annotations__',
  '__class__',
  '__delattr__',
  '__dict__',
  '__dir__',
  '__doc__',
  '__eq__',
  '__format__',
  '__ge__',
  '__getattr__',
  '__getattribute__',
  '__gt__',
  '__hash__',
  '__init__',
  '__init_subclass__',
  '__le__',
  '__lt__',
  '__module__',
  '__ne__',
  '__new__',
  '__reduce__',
  '__reduce_ex__',
  '__repr__',
  '__setattr__',
  '__sizeof__',
  '__str__',
  '__subclasshook__',
  '__weakref__']
```

```
'_get_or_create',
'_inferSchema',
'_instantiatedContext',
'_jsc',
'_jsqlContext',
'_jvm',
'_sc',
'_ssql_ctx',
'cacheTable',
'clearCache',
'createDataFrame',
'createExternalTable',
'dropTempTable',
'getConf',
'getOrCreate',
'newSession',
'range',
'read',
'readStream',
'reregisterDataFrameAsTable',
'reregisterFunction',
'reregisterJavaFunction',
'setConf',
'sparkSession',
'sql',
'streams',
'table',
'tableNames',
'tables',
'udf',
'udtf',
```